



```
In [1]: #Import libraries
import pandas as pd
import numpy as np
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
```

```
In [3]: #load data
saldf=pd.read_csv('/content/Salary Data.csv')
```

```
In [4]: #view data
saldf.head()
```

Out[4]:

	Age	Gender	Education Level	Job Title	Years of Experience	Salary
0	32.0	Male	Bachelor's	Software Engineer	5.0	90000.0
1	28.0	Female	Master's	Data Analyst	3.0	65000.0
2	45.0	Male	PhD	Senior Manager	15.0	150000.0
3	36.0	Female	Bachelor's	Sales Associate	7.0	60000.0
4	52.0	Male	Master's	Director	20.0	200000.0

```
In [5]: saldf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 375 entries, 0 to 374
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Age              373 non-null    float64
 1   Gender           373 non-null    object  
 2   Education Level  373 non-null    object  
 3   Job Title        373 non-null    object  
 4   Years of Experience 373 non-null    float64
 5   Salary           373 non-null    float64
dtypes: float64(3), object(3)
memory usage: 17.7+ KB
```

```
In [6]: saldf.isnull().sum()
```

```
Out[6]:
```

	0
Age	2
Gender	2
Education Level	2
Job Title	2
Years of Experience	2
Salary	2

dtype: int64

```
In [7]: #prepare independent and dependent variables  
inp=saldf[['Years of Experience']]  
out=saldf['Salary']
```

```
In [8]: #create an instance of LR  
LR=LinearRegression()
```

```
In [9]: #train the model  
df_cleaned = saldf.dropna(subset=['Years of Experience', 'Salary'])  
  
inp_cleaned = df_cleaned[['Years of Experience']]  
out_cleaned = df_cleaned['Salary']  
  
LR.fit(inp_cleaned, out_cleaned)
```

```
Out[9]: ▾ LinearRegression ⓘ ⓘ  
LinearRegression()
```

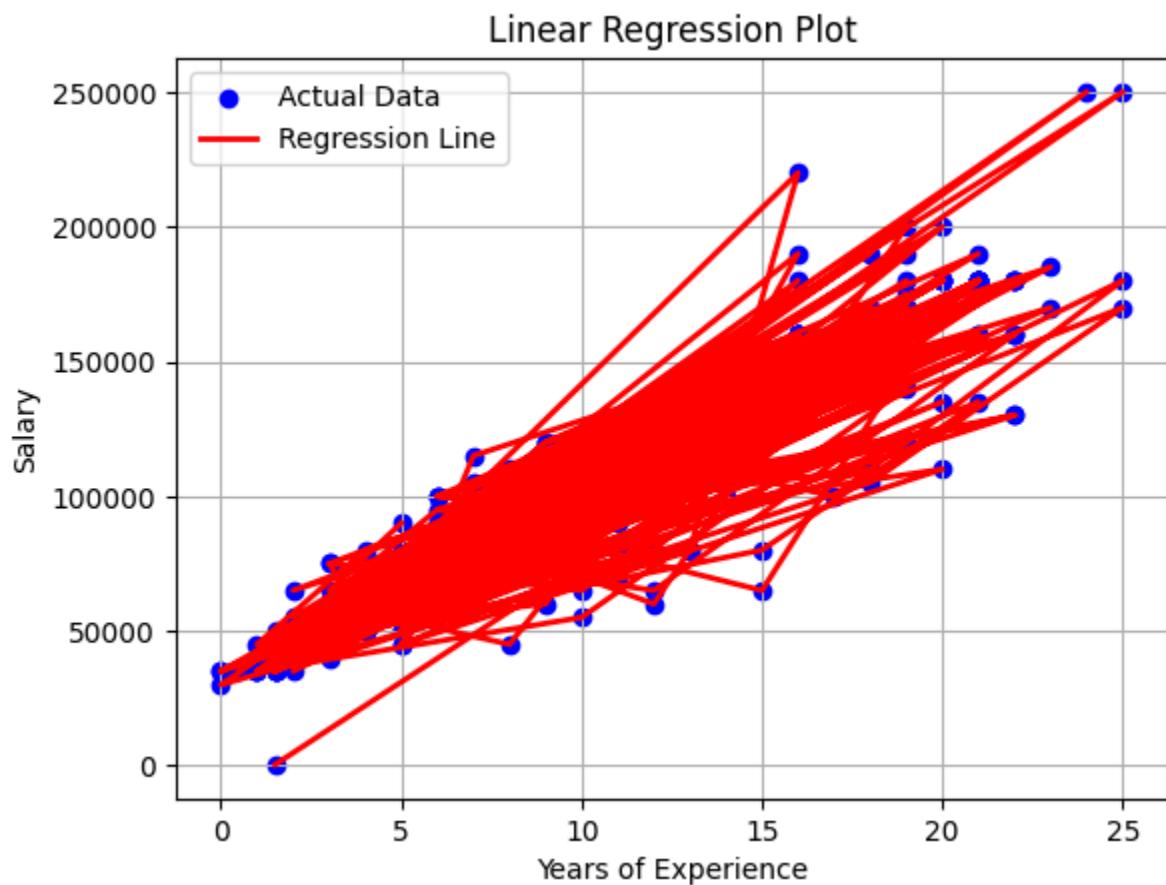
```
In [10]: #predict when experience is 5years  
LR.predict([[5]])
```

```
/usr/local/lib/python3.12/dist-packages/sklearn/utils/validation.py:2739: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names  
warnings.warn(
```

```
Out[10]: array([66143.76948947])
```

```
In [11]: plt.scatter(inp,out, color='blue', label='Actual Data')  
plt.plot(inp,out, color='red', linewidth=2, label='Regression Line')  
  
plt.xlabel('Years of Experience')  
plt.ylabel('Salary')  
plt.title('Linear Regression Plot')  
plt.legend()
```

```
plt.grid(True)  
plt.show()
```



```
In [12]: import pickle  
with open("Salaryprediction.pkl", "wb") as f:  
    pickle.dump(LR, f)  
from google.colab import files  
files.download('Salaryprediction.pkl')
```